

Agenda

- **Convolutional & pooling layers**
- **Convolutional neural networks**
- **Feature visualization**
- **Applications**

Images as 2D Arrays

- Grayscale image is a 2D array of pixel values
- Color images are 3D array
 - 3rd dimension is color (e.g., RGB)
 - Called “channels”




0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0

Convolution Neural Network: CNN

We know it is good to learn a small model.

From this fully connected model, do we really need all the edges?


Can some of these be shared?



Consider learning an image:

Some patterns are much smaller than the whole image


Can represent a small region with fewer parameters



Same pattern appears in different places:


They can be compressed!

What about training a lot of such “small” detectors
and each detector must “move around”.



A convolutional layer

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.




What Is Translation?

A **translation** is a geometric transformation that shifts all points in a given direction and by the same distance. Alternatively, it can be interpreted as sliding the origin of the coordinate system by the same amount but in the opposite direction.

The translation can be expressed mathematically as the vector sum of a constant vector \mathbf{v} to each point \mathbf{x} :

$$T_{\mathbf{v}}(\mathbf{x}) = \mathbf{x} + \mathbf{v}$$



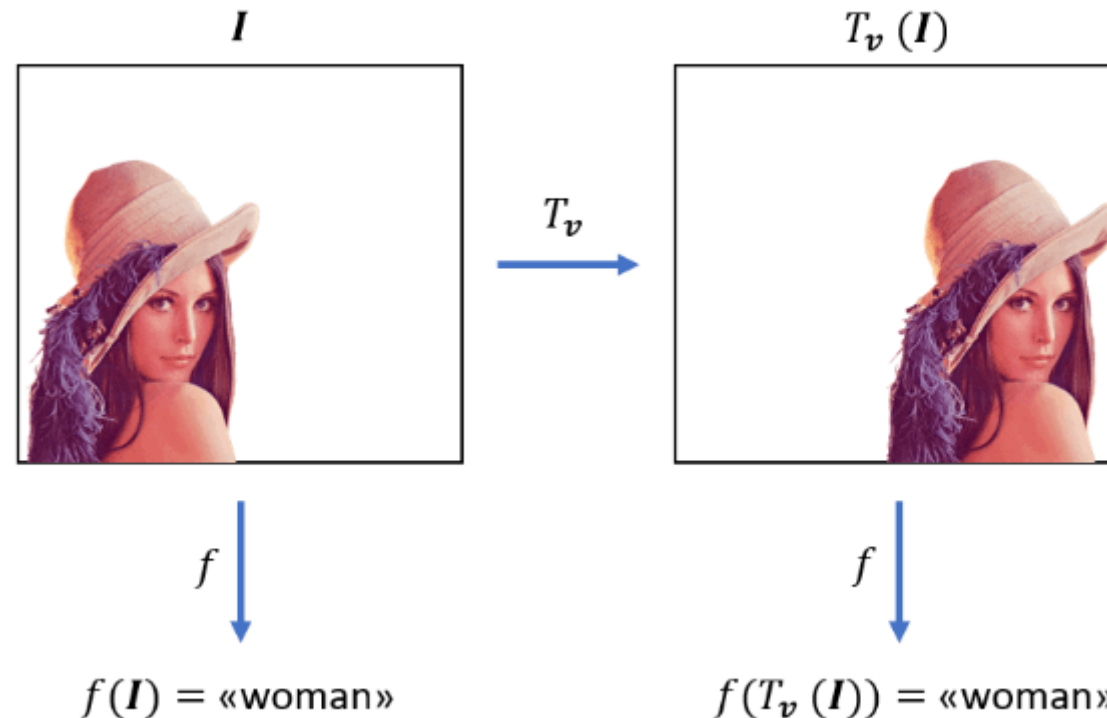
(a)



(b)

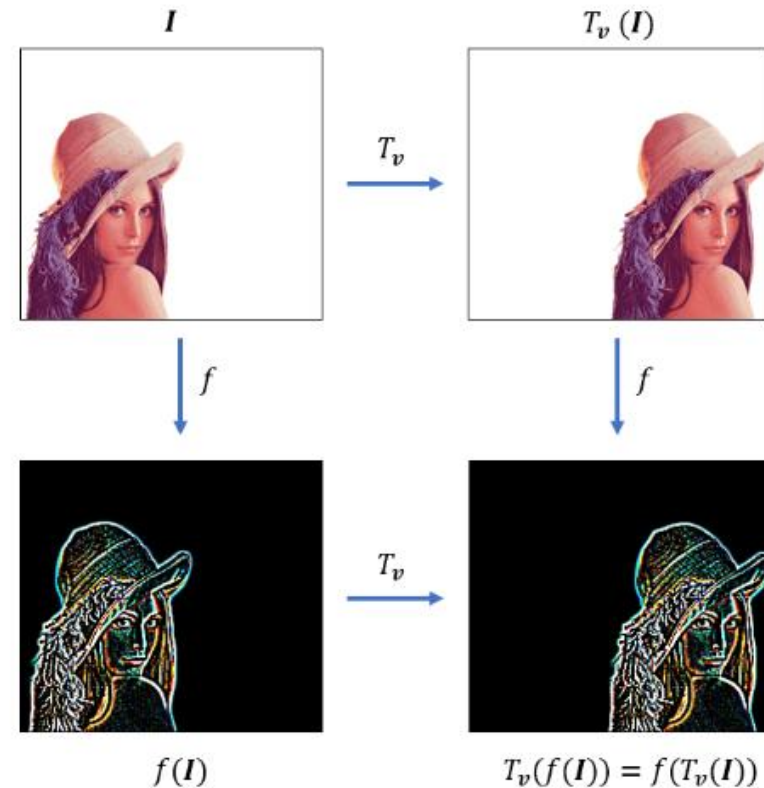
Translation Invariance

- **Translation invariance** - is a property of a system or model where its performance or output does not change when the input data is translated or shifted.



Translation Equivariance

Translation equivariance - is a property of a function or an operation where the output's structure does not change even if the input is translated.




A function f is said equivariant to a function g if and only if:

$$f(g(x)) = g(f(x))$$


In other words, f is equivariant to g if the order of application does not change the result of the composite function.

Structure in Images

- Use layers that capture structure



Convolution layers
(Capture equivariance)



Pooling layers
(Capture invariance)

Vertical edge detection


10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*



1	0	-1
1	0	-1
1	0	-1

=


0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



*



Convolution Filters



$$\text{output}[i, j] = \sum_{\tau=0}^{k-1} \sum_{\gamma=0}^{k-1} \text{filter}[\tau, \gamma] \cdot \text{image}[i + \tau, j + \gamma]$$

2D Convolution Filters

- **Given:**

- A 2D input x
- A 2D $h \times w$ kernel k

- The 2D convolution is:

$$y[s, t] = \sum_{\tau=0}^{h-1} \sum_{\gamma=0}^{w-1} k[\tau, \gamma] \cdot x[s + \tau, t + \gamma]$$

2D Convolution Filters

-1	-1	-1
2	2	2
-1	-1	-1

Horizontal lines

-1	2	-1
-1	2	-1
-1	2	-1


Vertical lines

-1	-1	2
-1	2	-1
2	-1	-1

45 degree lines

2	-1	-1
-1	2	-1
-1	-1	2

135 degree lines



Example Edge Detection Kernels

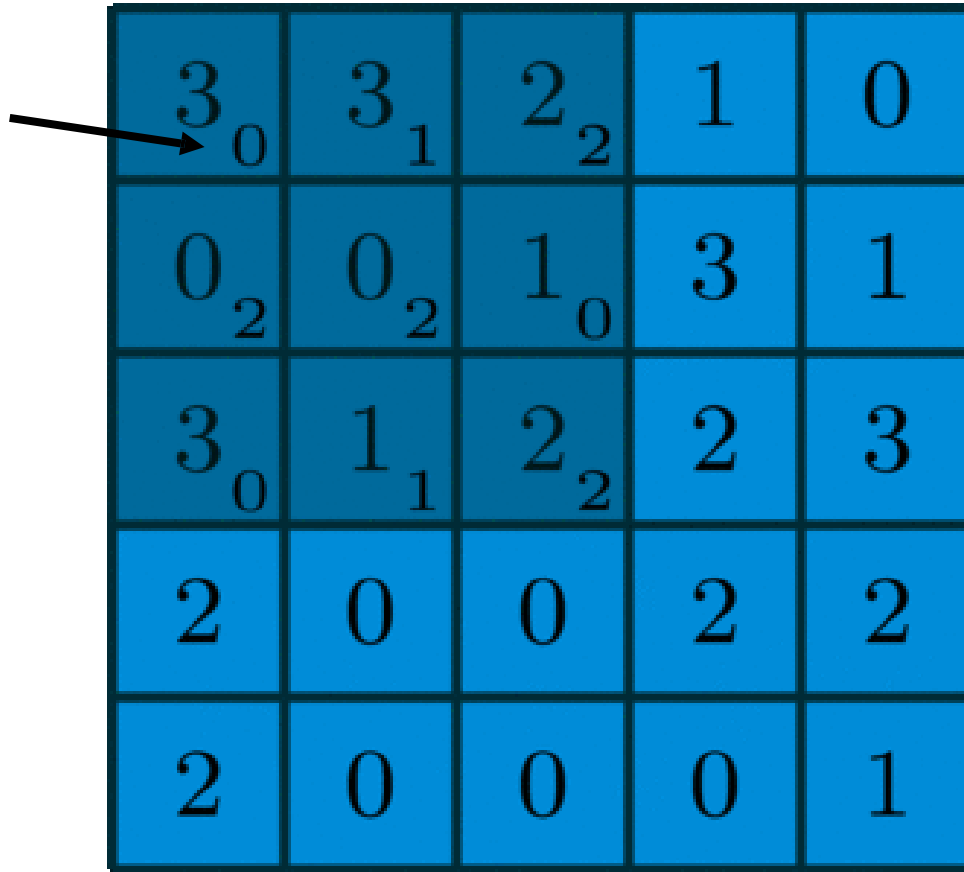
Result of Convolution with Horizontal Kernel

2D Convolution Filters

- Historically (until late 1980s), kernel parameters were handcrafted
 - E.g., “edge detectors”
- In convolutional neural networks, they are learned
 - Essentially a linear layer with fewer “connections”
 - Backpropagate as usual!

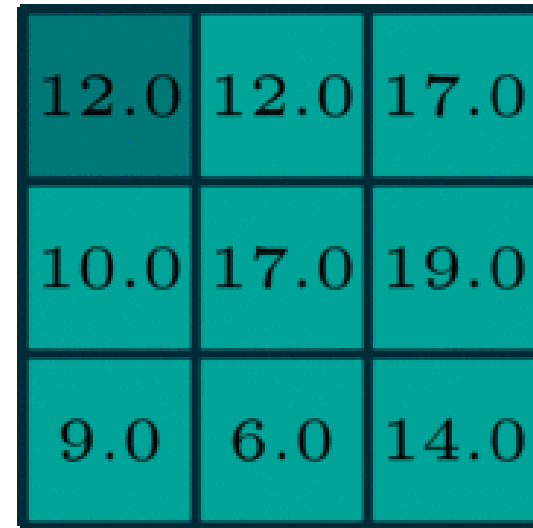
Convolution Layers

Learnable
parameters



A 5x5 grid representing a convolution kernel. The first three columns are highlighted in dark blue, indicating learnable parameters, while the last two columns are light blue, indicating fixed values. An arrow points from the text 'Learnable parameters' to the top-left cell of the dark blue region.

3_0	3_1	2_2	1	0
0_2	0_2	1_0	3	1
3_0	1_1	2_2	2	3
2	0	0	2	2
2	0	0	0	1

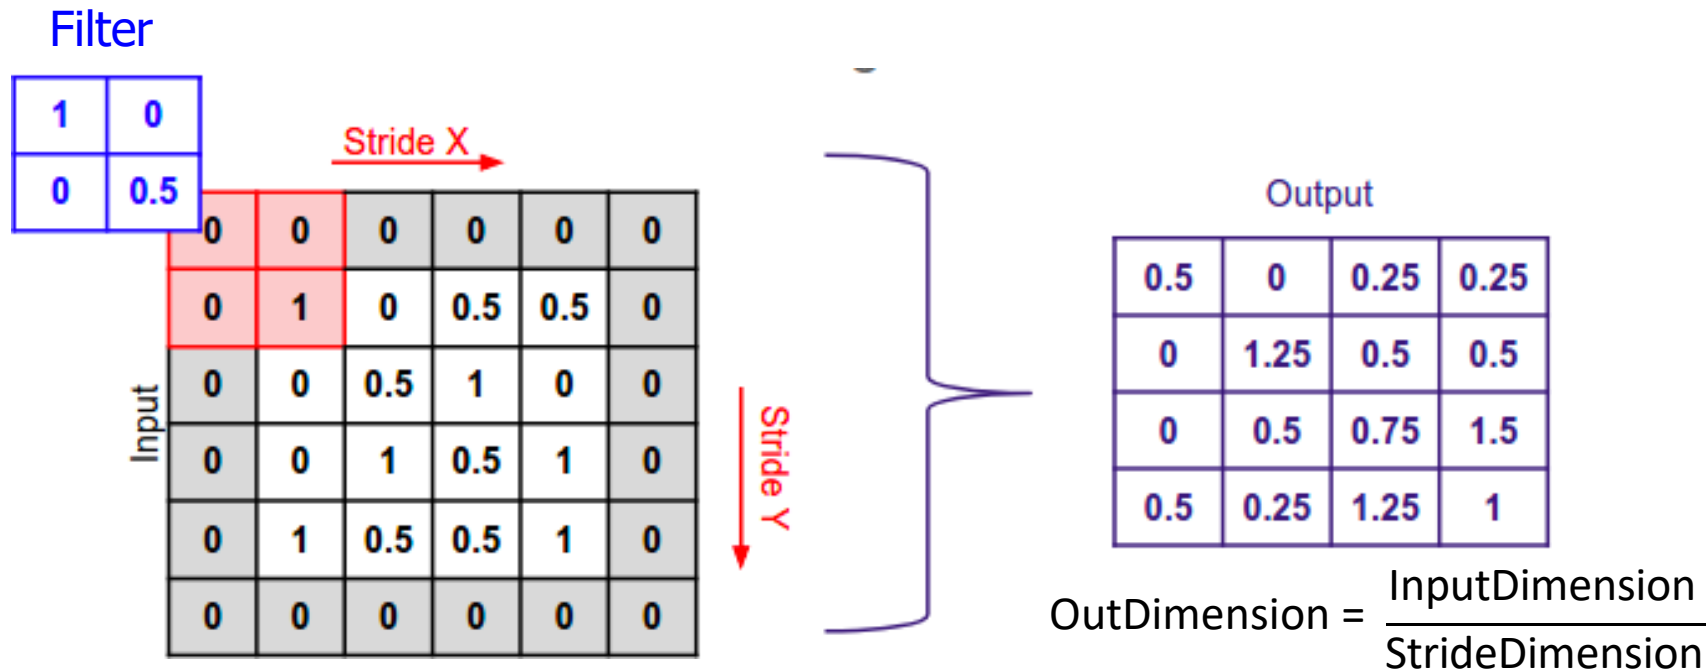


A 3x3 grid representing the output feature map, with values ranging from 6.0 to 19.0. The cells are colored in shades of teal.

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

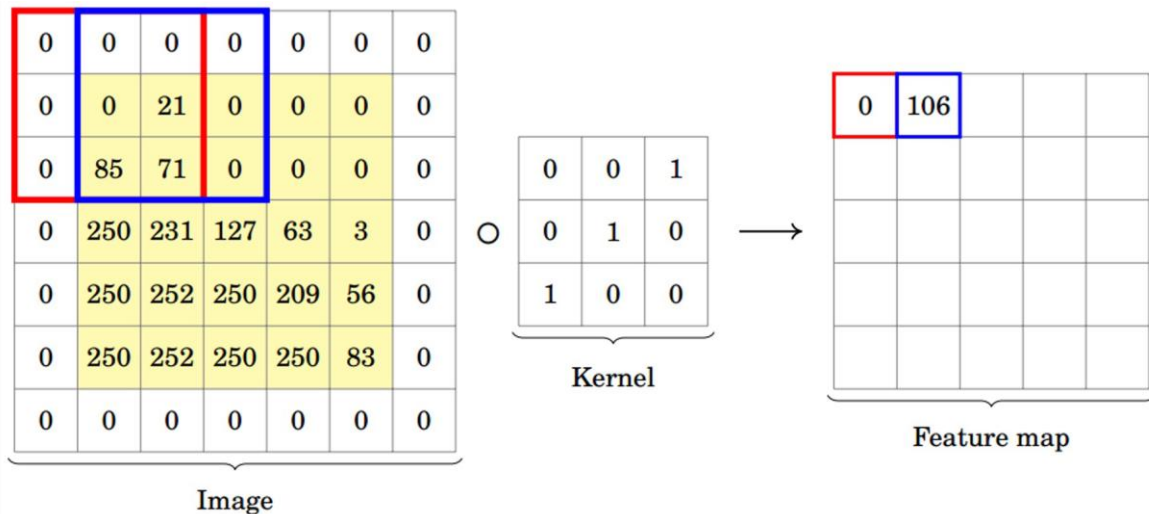
Convolution Layer Parameters

- **Stride:** How many pixels to skip (if any)
 - **Default:** Stride of 1 (no skipping)




Convolution Layer Parameters

- **Padding:** Add zeros to edges of image to capture ends
 - **Default:** No padding



stride = 1, zero-padding = 1




stride = 2, zero-padding = 1

Convolution Layer Parameters


- **Summary:** Hyperparameters
 - Kernel size
 - Stride
 - Amount of zero-padding
 - Output channels
- Together, these determine the relationship between the input tensor shape and the output tensor shape
- Typically, also use a single bias term for each convolution filter

Convolution Layers




Example


- Kernel size 3, stride 2, padding 1
- 3 input channels
 - Hence kernel size $3 \times 3 \times 3$
- 2 output channels
 - Hence 2 kernels
- Total # of parameters:
 - $(3 \times 3 \times 3 + 1) \times 2 = 56$



Pooling Layers




Pooling Layers




$$\text{output}[0,0] = \max_{0 \leq \tau < k} \max_{0 \leq \gamma < k} \text{image}[0 + \tau, 0 + \gamma]$$

Pooling Layers




$$\text{output}[0,1] = \max_{0 \leq \tau < k} \max_{0 \leq \gamma < k} \text{image}[0 + \tau, 1 + \gamma]$$

Pooling Layers



$$\text{output}[0,2] = \max_{0 \leq \tau < k} \max_{0 \leq \gamma < k} \text{image}[0 + \tau, 2 + \gamma]$$

Pooling Layers



$$\text{output}[i, j] = \max_{0 \leq \tau < k} \max_{0 \leq \gamma < k} \text{image}[i + \tau, j + \gamma]$$

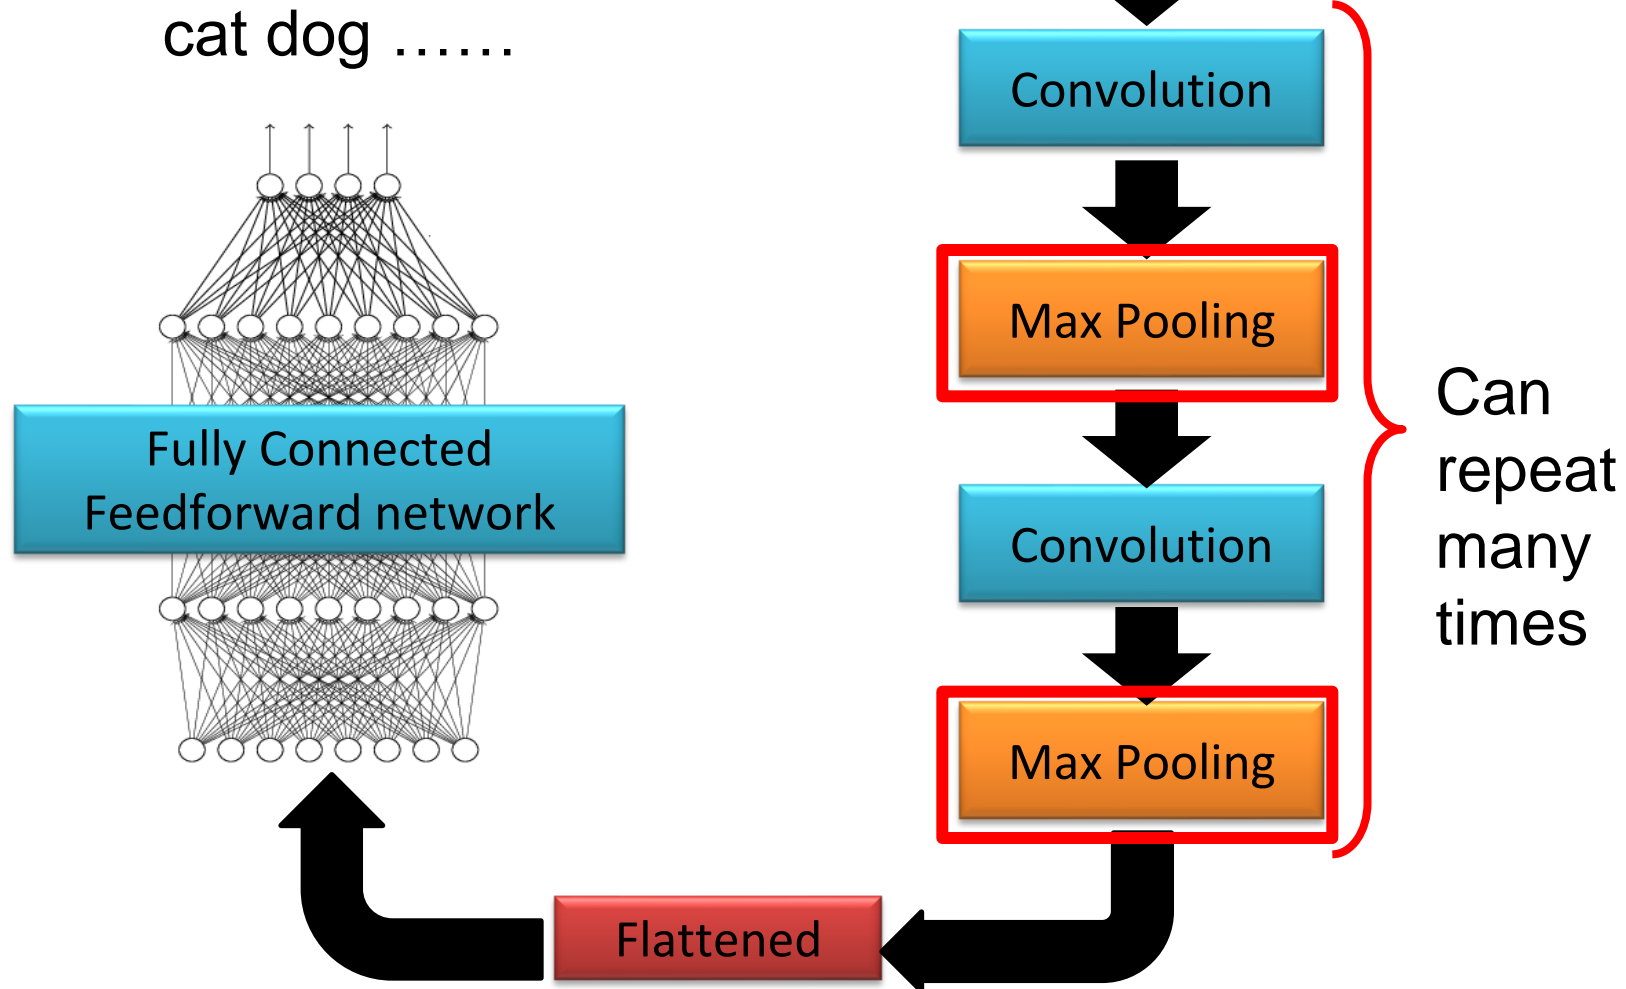
Pooling Layers

- **Summary:** Hyperparameters
 - Kernel size
 - Stride (usually >1)
 - Amount of zero-padding
 - Pooling function (almost always “max”)
- Together, these determine the relationship between the input tensor shape and the output tensor shape
- **Note:** Unlike convolution, pooling operates on channels separately
 - Thus, n input channels $\rightarrow n$ output channels

Summary: Convolution vs. Pooling

- **Convolution layers:** Translation equivariant
 - If object is translated, convolution output is translated by same amount
 - Produce “image-shaped” features that retain associations with input pixels
- **Pooling layers:** Translation invariant
 - Binning to make outputs insensitive to translation
 - Also reduces dimensionality
- Combined in modern architectures
 - Convolution to construct equivariant features
 - Pooling to enable invariance

The whole CNN




Example

- Suppose we want to predict whether an image depicts Cartesian axes

$$\begin{bmatrix} 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & -2 & 0 & -2 \end{bmatrix}$$

input image



target (binary) label

Example

- **Step 1:** Convolve the image with two filters
 - No padding, stride 1
- **Step 2:** Run max pooling

$$\begin{bmatrix} -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \end{bmatrix}, \begin{bmatrix} -\frac{1}{2} & -\frac{1}{2} & -\frac{1}{2} \\ 1 & 1 & 1 \\ -\frac{1}{2} & -\frac{1}{2} & -\frac{1}{2} \end{bmatrix}$$

convolution filters

Example

$$\begin{bmatrix} 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & -2 & 0 & -2 \end{bmatrix} \quad \begin{bmatrix} -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \end{bmatrix}$$

Example

$$\begin{bmatrix} 0 & 1 & 0 \\ 0 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \end{bmatrix}$$



$$\begin{bmatrix} 2 & . \\ . & . \end{bmatrix}$$

$$\begin{aligned} &\left(0 \times \frac{-1}{2}\right) + (1 \times 1) + \left(0 \times \frac{-1}{2}\right) \\ &\left(0 \times \frac{-1}{2}\right) + (1 \times 1) + \left(0 \times \frac{-1}{2}\right) \\ &\left(1 \times \frac{-1}{2}\right) + (1 \times 1) + \left(1 \times \frac{-1}{2}\right) = 2 \end{aligned}$$

Example

$$\begin{bmatrix} 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & 1 & 0 & 1 \end{bmatrix} \begin{bmatrix} -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \end{bmatrix} \rightarrow \begin{bmatrix} 2 & -2 \\ \cdot & \cdot \end{bmatrix}$$

Example

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & -2 & 0 \end{bmatrix}$$

$$\begin{bmatrix} -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \end{bmatrix}$$



$$\begin{bmatrix} 2 & -2 \\ -1 & . \end{bmatrix}$$

Example

$$\begin{bmatrix} 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & -2 & 0 & -2 \end{bmatrix} \begin{bmatrix} -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \end{bmatrix} \rightarrow \begin{bmatrix} 2 & -2 \\ -1 & 1 \end{bmatrix}$$

Example

The diagram illustrates a convolution operation. On the left, an input image is represented as a 4x4 matrix. Two arrows branch from this matrix to two separate 3x3 convolution kernels in the middle. Each kernel is then convolved with the input image, resulting in a 2x2 feature map on the right. The top path uses a kernel with values $-\frac{1}{2}$, 1, and $-\frac{1}{2}$ to produce a feature map with values 2 and -2 in the first row, and -1 and 1 in the second row. The bottom path uses a kernel with values $-\frac{1}{2}$, $-\frac{1}{2}$, and $-\frac{1}{2}$ in the first row, 1, 1, and 1 in the second row, and $-\frac{1}{2}$, $-\frac{1}{2}$, and $-\frac{1}{2}$ in the third row, to produce a feature map with values -1 and $-\frac{1}{2}$ in the first row, and $\frac{7}{2}$ and 4 in the second row.

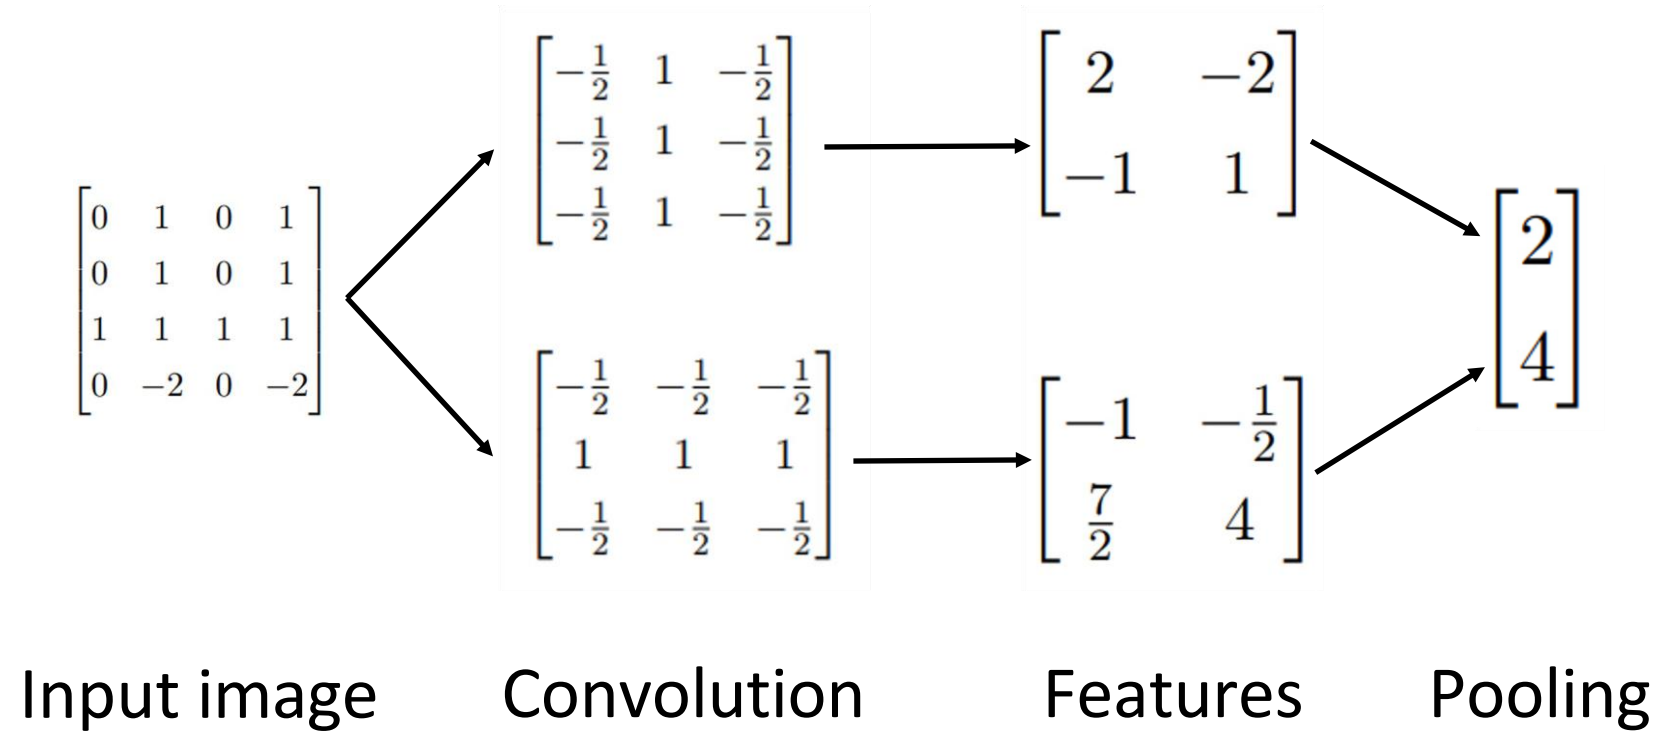
$$\begin{bmatrix} 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 1 \\ 0 & -2 & 0 & -2 \end{bmatrix} \begin{matrix} \nearrow \\ \searrow \end{matrix}$$
$$\begin{bmatrix} -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \\ -\frac{1}{2} & 1 & -\frac{1}{2} \end{bmatrix} \longrightarrow \begin{bmatrix} 2 & -2 \\ -1 & 1 \end{bmatrix}$$
$$\begin{bmatrix} -\frac{1}{2} & -\frac{1}{2} & -\frac{1}{2} \\ 1 & 1 & 1 \\ -\frac{1}{2} & -\frac{1}{2} & -\frac{1}{2} \end{bmatrix} \longrightarrow \begin{bmatrix} -1 & -\frac{1}{2} \\ \frac{7}{2} & 4 \end{bmatrix}$$

Input image

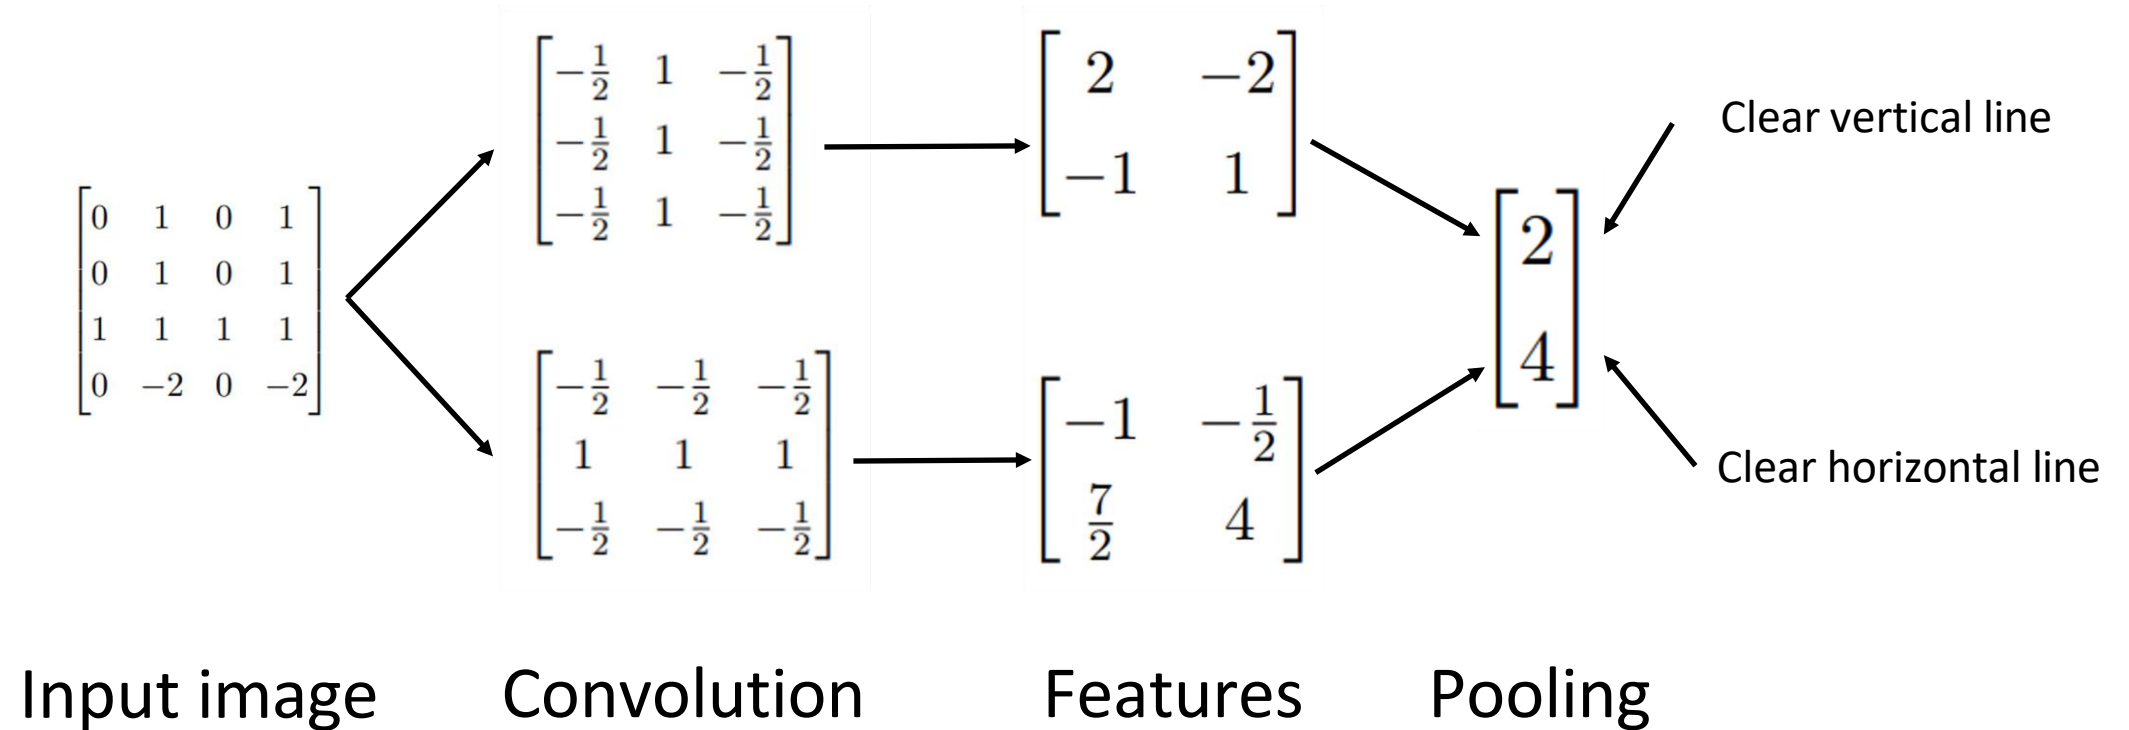
Convolution

Features

Example



Example



$n \times n$ image $f \times f$ filter

padding p stride s

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$

Agenda

- **Convolutional & pooling layers**
- **Convolutional neural networks**
- **Feature visualization**
- **Applications**

Example Architecture: AlexNet

- **ImageNet dataset**


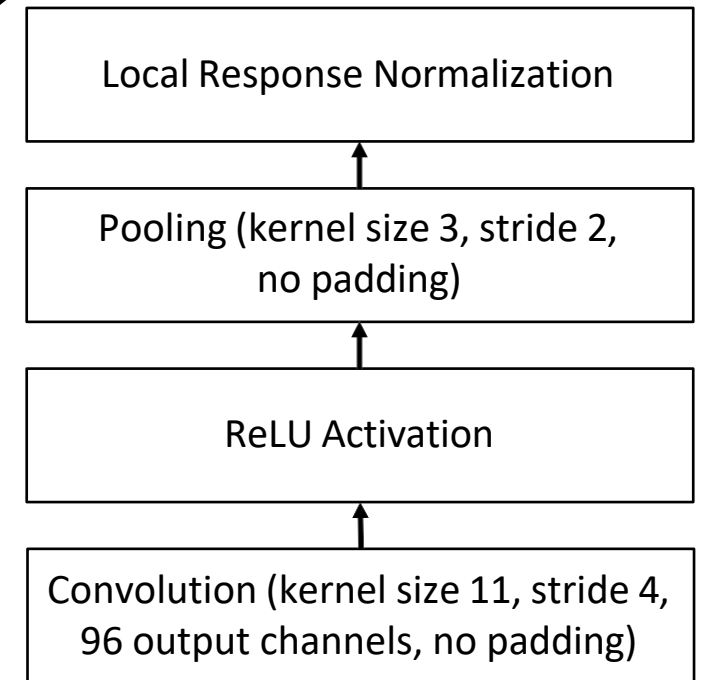
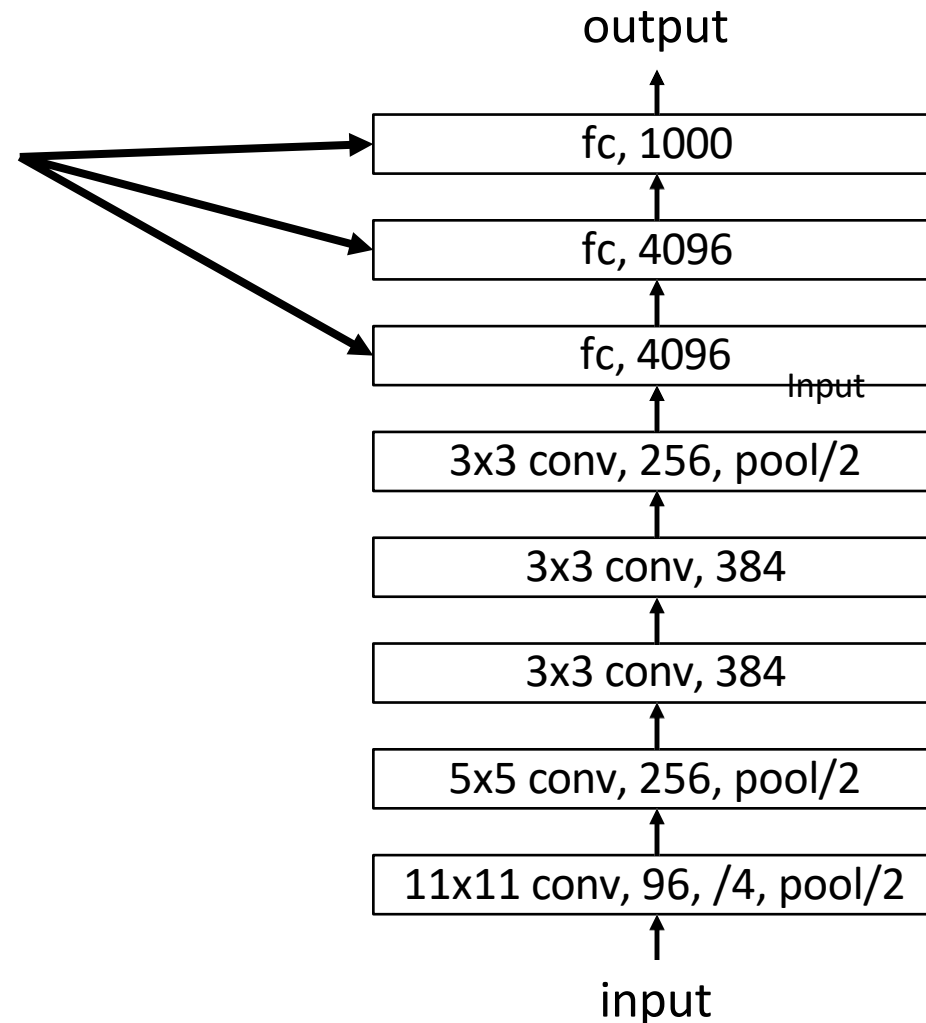
- 1000 class image classification problem (e.g., grey fox, tabby cat, barber chair)
- >1M image-label pairs gathered from internet and crowdsourced labels

- **AlexNet Architecture (Krizhevsky 2012)**


- Historically important architecture
- Image classification network (~60M parameters)
- Trained using GPUs on ImageNet dataset
- Huge improvement in performance compared to prior state-of-the-art

Example Architecture: AlexNet

Fully connected
(i.e., linear) layers




Example Architecture: AlexNet




Aside: Local Response Normalization

- Highlights areas where the feature maps change
- Historically a standard layer, but no longer used
- Also called “contrastive normalization”




Convolutional Neural Networks

- “Convolutional layer” often refers to sequence of layers
- **Modern sequence of layers**
 - Convolution -> Batch Normalization -> Pooling -> ReLU
 - Convolution -> Batch Normalization -> ReLU -> Pooling
- Can also omit pooling (especially for very deep neural networks)



<https://towardsdatascience.com/batch-norm-explained-visually-how-it-works-and-why-neural-networks-need-it-b18919692739>

Mini-batch: Activations



Batch Norm

Mean and Std Dev

$$\mu_i = \frac{1}{M} \sum A_i$$
$$\sigma_i = \sqrt{\frac{1}{M} \sum (A_i - \mu)^2}$$

2

Normalize

$$\hat{A}_i = \frac{A_i - \mu_i}{\sigma_i}$$

3

Scale and Shift

$$B\tilde{N}_i = \gamma \odot \hat{A}_i + \beta$$

4

Beta
(β)

Gamma
(γ)

Output
(bn)

Moving Average

$$\mu_{mov_i} = \alpha \mu_{mov_i} + (1 - \alpha) \mu_i$$

$$\sigma_{mov_i} = \alpha \sigma_{mov_i} + (1 - \alpha) \sigma_i$$



Moving Avg
(mean)

Moving Avg
(Var)

5


Residual Connections

- **Challenges with deeper networks**
 - Overfitting?
 - No, 56 layer network underfits!



Residual Connections

- **Challenges with deep networks**
 - Overfitting?
 - No, 56 layer network underfits!
- **Optimization/representation**
 - Difficulty representing the identity function!
- **Solution:** “Skip” connections
 - Facilitate direct feedback from loss
 - Easy to represent identity function



Residual Connections

- **Residual layers:** Given any convolutional layer $F(x)$, use


$$H(x) = F(x) + x$$

- Two views of residual connections:
 - **View 1:** Providing shortcuts to gradients on the backward pass
 - **View 2:** Allow each “residual block” to fit the residual error (boosting!)

$$F(x) = H(x) - x$$

Residual Networks

- Stack lots of residual blocks!
 - Kernel size 3, no padding, stride 1, no pooling
 - Reduce feature dimensions by using stride 2 once every K blocks
 - Maintains feature size to build very deep nets



Residual Networks

- For deeper networks, improve efficiency through 1x1 convolutions
- Many other improvements since 2015!
 - E.g., “ResNeXt”, “Identity Mappings”, “ConvNeXt” etc.





Figure 5. A deeper residual function \mathcal{F} for ImageNet. Left: a building block (on 56×56 feature maps) as in Fig. 3 for ResNet-34. Right: a “bottleneck” building block for ResNet-50/101/152.

Evolution of Neural Networks



Evolution of Neural Networks

AlexNet, 8 layers
(ILSVRC 2012)
~60M params



VGG, 19 layers
(ILSVRC 2014)
~140M params



ResNet, 152 layers
(ILSVRC 2015)
Less computation
in forward pass
than VGGNet!
Back to 60M params



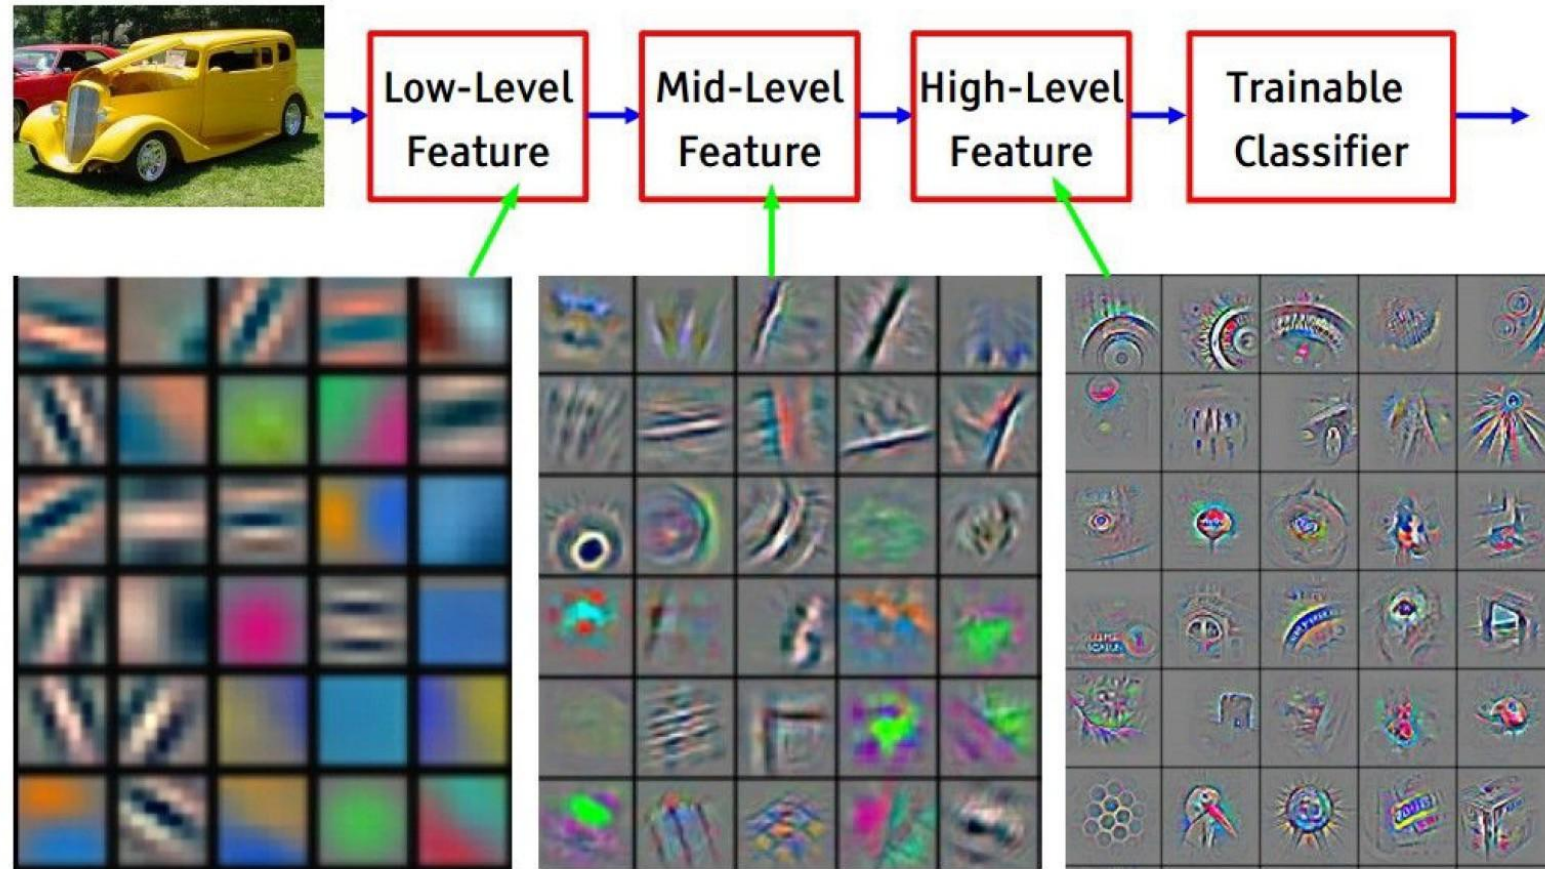
GoogleNet, 22 layers
(ILSVRC 2014)
~5M params



Agenda


- **Convolutional & pooling layers**
- **Convolutional neural networks**
- **Feature visualization**
- **Applications**

Feature Visualization




Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]


Layer 1




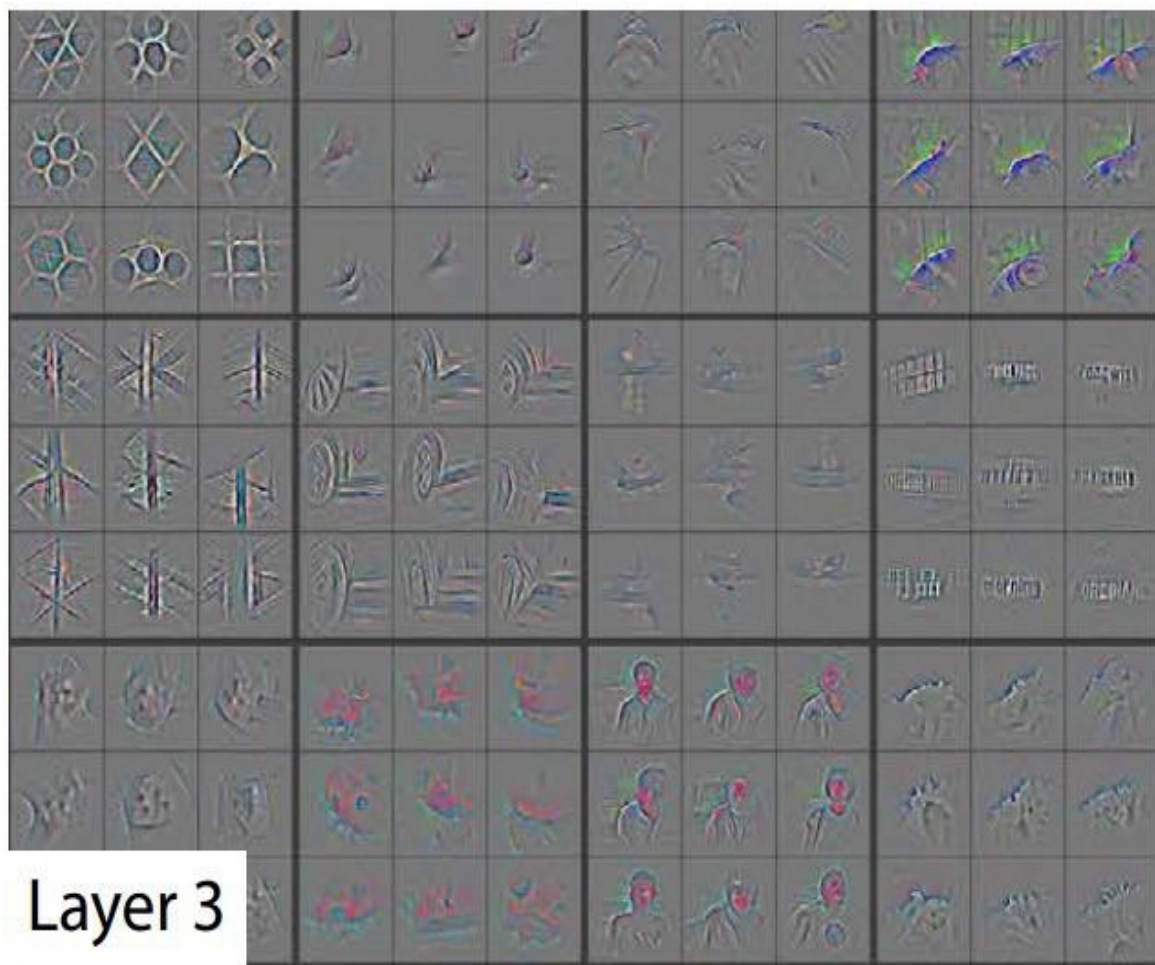
Layer 1



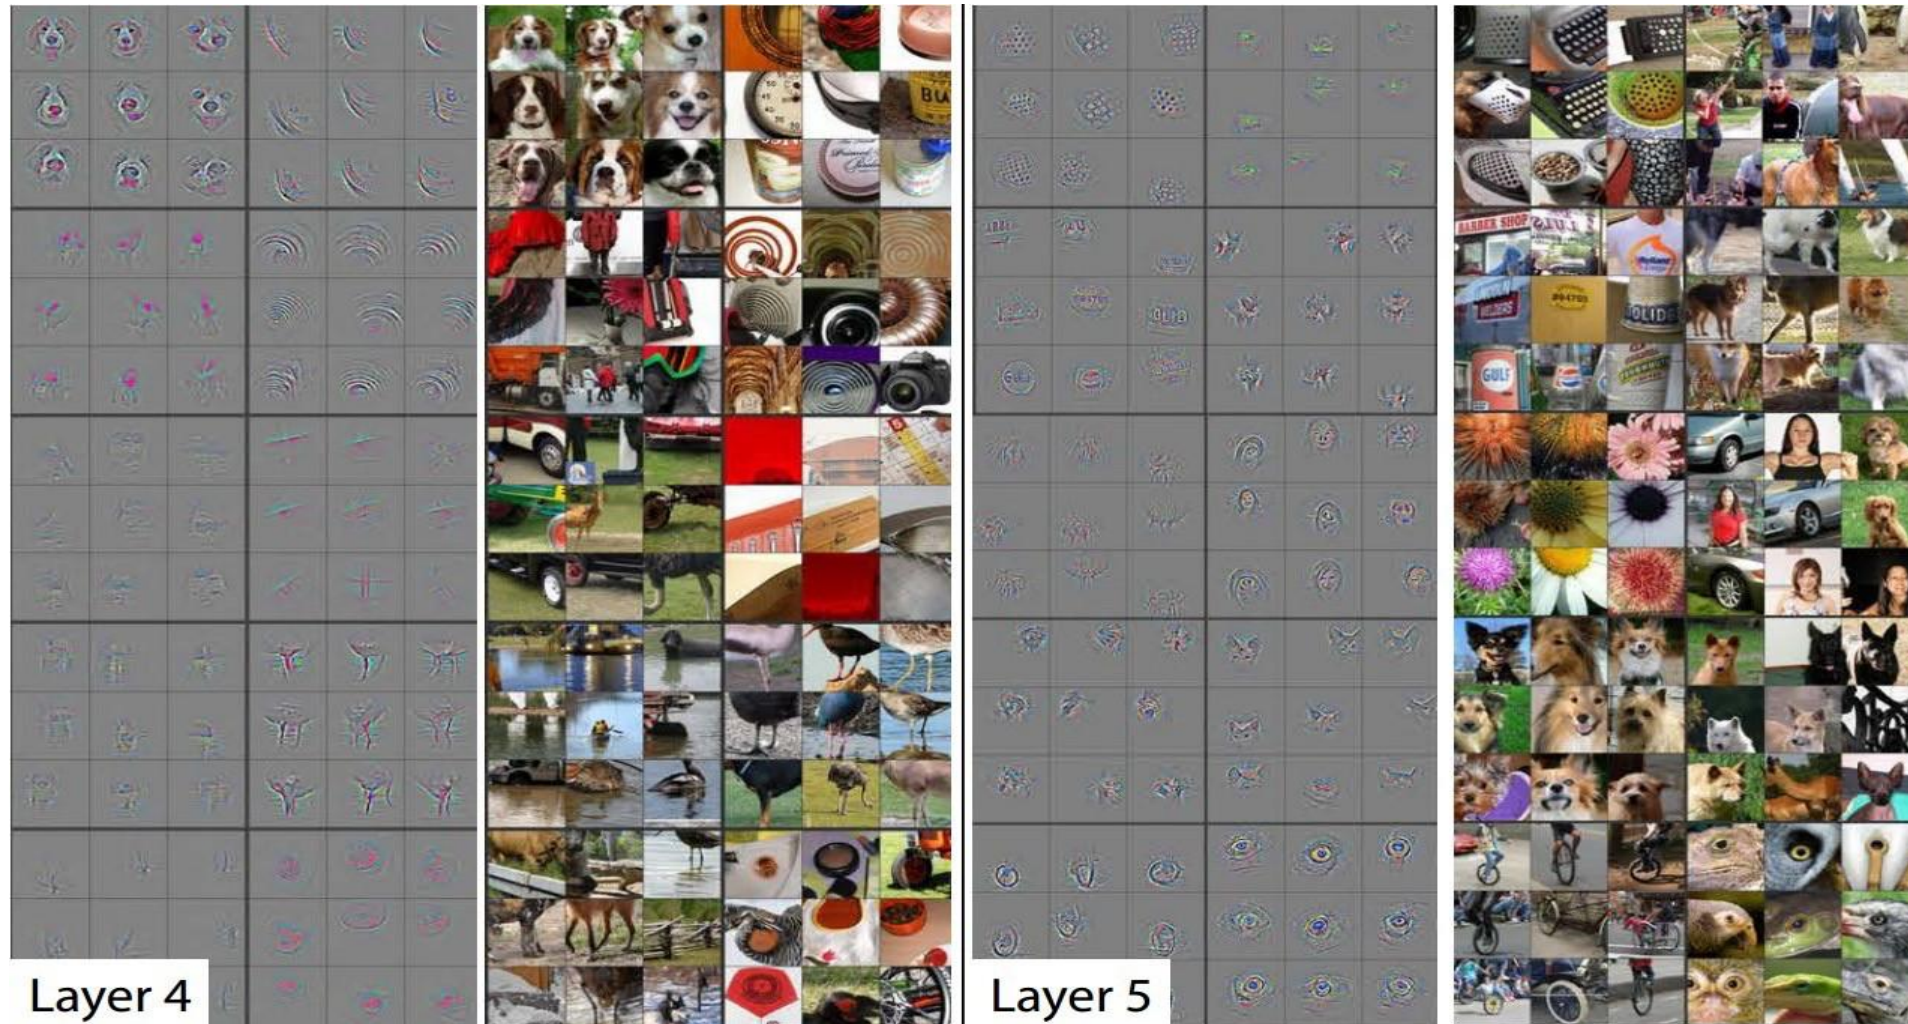
Layer 2



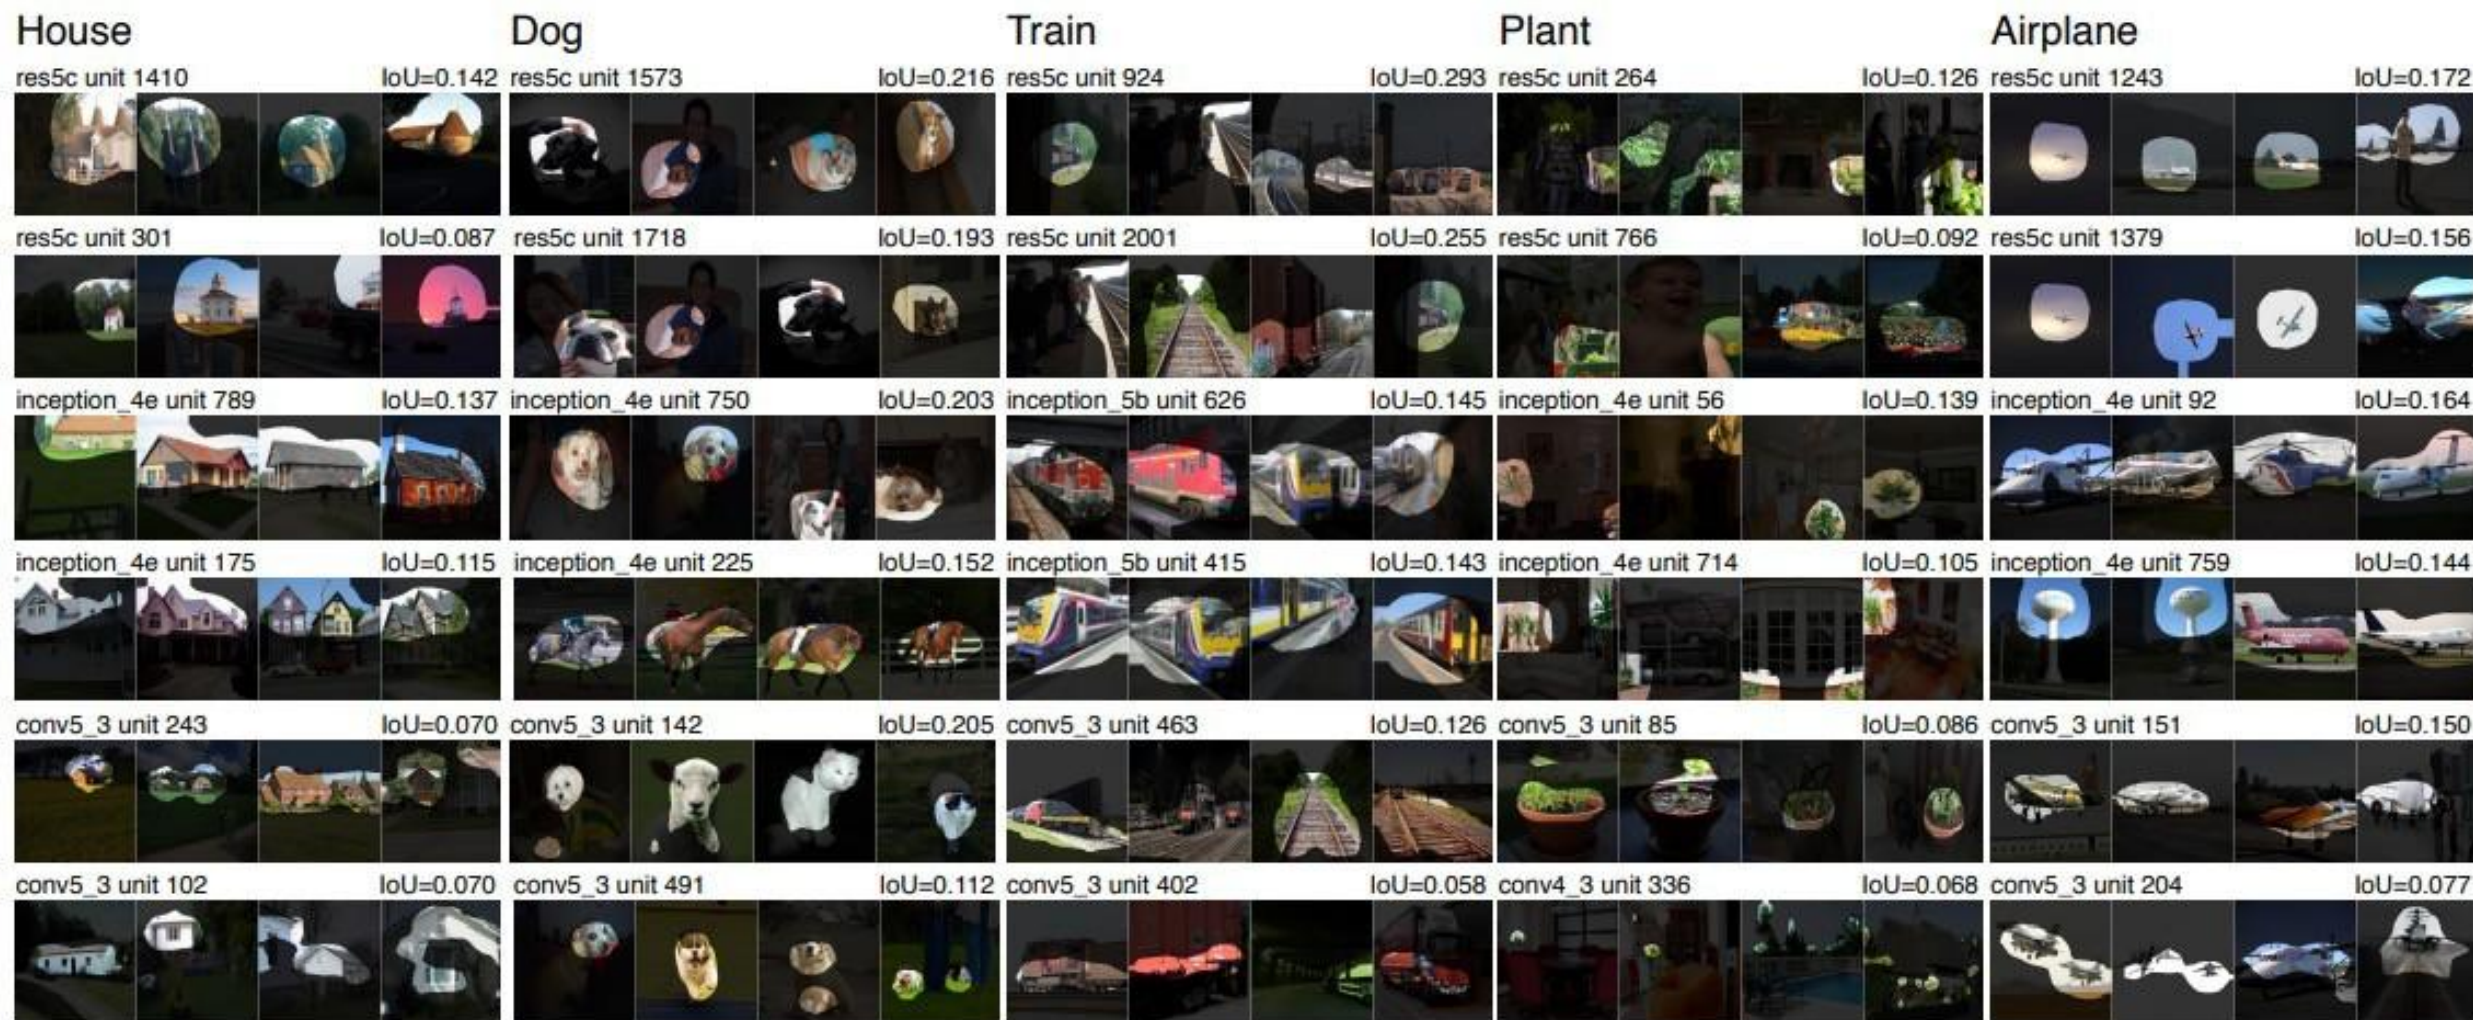
Layer 3



Layer 3



Neural Network Dissection



What About Small Datasets?

- **Transfer learning:** We can reuse trained concepts!
 - Since CNNs trained on ImageNet appear to learn general features
 - We can reuse these models in some way to perform new tasks
- **Strategy 1:** Feature extraction
 - Remove final (softmax) layer and replace with a new one
 - Train only the new layer
- **Strategy 2:** Finetuning
 - Do the same thing but train end-to-end

What About Small Datasets?

- **New dataset is similar to the original dataset**
 - Can use very small datasets
 - Both strategies work
- **New dataset is different from original dataset**
 - Transfer learning still works!
 - Moderate-sized datasets
 - Finetune end-to-end
 - **Examples:** Medical images, audio spectrograms, etc.

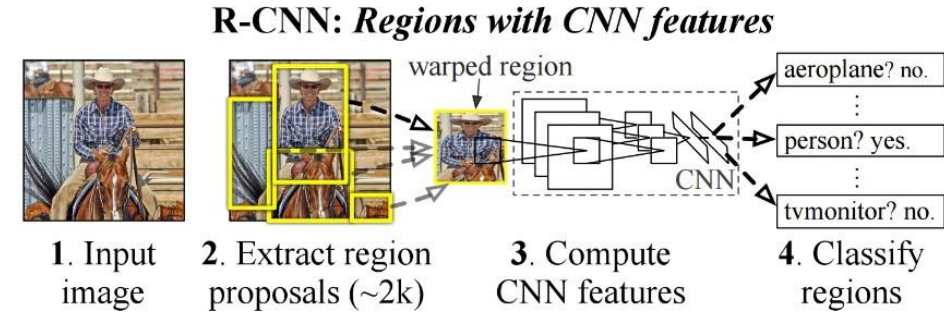
Agenda

- **Convolutional & pooling layers**
- **Convolutional neural networks**
- **Feature visualization**
- **Applications**

Applications

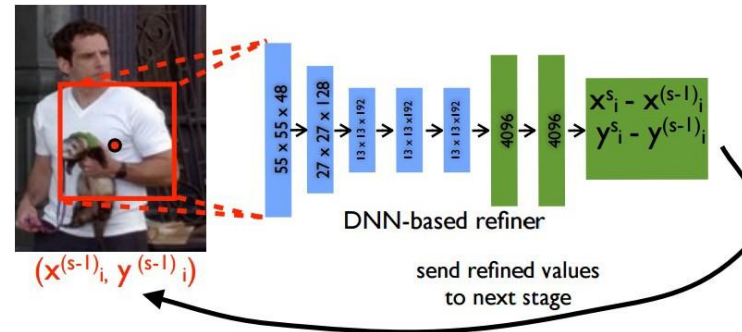
Object detection

[Girshick et al. CVPR14]




Pose detection (regression)

[Toshev et al. CVPR14]



Semantic segmentation

[Long et al. CVPR15]



Applications

Similarity metric learning






Image generation




Low-level image processing: (superresolution, deblurring, image quality etc.)



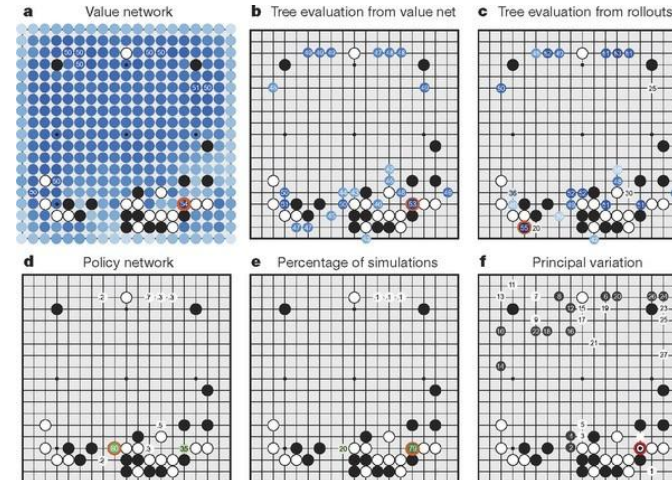
Applications: Game Playing

CNN + Reinforcement learning

[Mnih et al, Nature '15]



[Silver et al, Nature '16]



Applications: Art Generation




See if you can tell artist originals from machine style imitations at:
<http://turing.deepart.io/>

Paper: [Gatys et al, "Neural ... Style", arXiv '15](#)
Code (torch): <https://github.com/jcjohnson/neural-style>



Structure in Images

- **Translation invariance** - is a property of a system or model where its performance or output does not change when the input data is translated or shifted.
 - Consider image classification (e.g., labels are cat, dog, etc.)
 - **Invariance:** If we translate an image, it does not change the category label



Source: Ott et al., Learning in the machine: To share or not to share?

Structure in Images

- **Translation equivariance** - is a property of a function or an operation where the output's structure does not change even if the input is translated.
 - Consider object detection (e.g., find the position of the cat in an image)
 - **Equivariance:** If we translate an image, the the object is translated similarly

